BACKBOND

AGENT RISK RATINGS & INSURANCE

BRAND SYSTEM · SOURCE OF TRUTH

Brand Backbone

The single reference for what BackBond is, how it sounds, and how it looks. Everything we make derives from this document.

VERSION 1.6
JUNE 2026
INTERNAL · CONFIDENTIAL

GOVERNED BY FAL
Focus · Alignment · Linkage

What we are, in one sentence.

THE LOCKED FOCUS

BackBond rates and insures AI agents, so companies can put them to work with coverage behind every action they take.

THE CATEGORY - WHAT MAKES IT DIFFERENT

AI agent insurance. Every agent carries its own coverage.

Coverage lives at the agent and meters by the transaction, priced the way AI is billed: per agent, per run. That is the ownable mechanic. Traditional carriers price the company once a year; only we price the agent, by the action.

One Focus, every register. The category and investor story, the buyer pitch, and the partner hook are plain-English variants of this one promise, and a stranger must recover the same Focus from any of them. The partner entry hook stays "Guardrails lower the odds. We cover the loss." (the tagline of record).

THE FOCUS DOES NOT CLAIM

Not AI safety. Not an AI-governance platform. Not "responsible AI." Not agent evaluation. And not total risk transfer: coverage pays the covered loss up to the limit bought, it never removes all downside. **We rate, and we insure.** Drift toward any of these is a ship-blocking error.

02 · POSITIONING

The rating-and-insurance layer.

BackBond is the rating-and-insurance layer for AI agents. We score each agent's risk and put real capital behind the score - per agent, per run - so a company can put its agents to work and answer the one question that stalls every deal, "who is liable when this acts?", with a policy instead of a disclaimer.

THE CATEGORY WE ARE BUILDING

Agent Risk Ratings & Insurance. A third-party rating layer plus the capital behind it. It is **AI agent insurance**: every agent carries its own coverage, priced by the run, the way AI is billed. Traditional carriers price the company once a year; only we price the agent, by the action. Every prior platform shift created a rating layer that captured disproportionate value (FICO, S&P, SOC 2, BitSight). Agents are next, and there is one king position.

THE EVAL-VS-ACTION DEFENSE - LOCKED

The score's pillars will be attacked as "model eval with insurance bolted on." The answer is never a pillar comparison. **BackBond is not another eval tool; it is insurance.** Eval vendors ship a PDF. We put capital behind the score and pay claims on action-level losses in the logs. The differentiator is the bond, not the pillars.

WEDGE TODAY

The methodology - a credible, third-party number an enterprise can demand in procurement.

MOAT TOMORROW

Loss data. Because we price by the action, every transaction we cover compounds a proprietary dataset of scored agents tied to real outcomes. Late entrants cannot replicate it.

03 · THE NAME

One name, one system.

THE SYSTEM

Company - BackBond. The brand.

Instrument - a bond / bonded. The agent is backed by a real insurance policy that pays the claim when it fails. ‘Bonded’ is evocative, as in ‘bonded and insured,’ not a literal surety bond.

Metric - the BackBond Score. 0–100. Always "BackBond Score," never a bare "Bond Score."

Conversation word - bondable / bondability. The everyday verb-form.

HOW IT SOUNDS IN THE ROOM

"What's its Bond?"

"BackBond Score of 78, Standard band - bondable today, Preferred after two fixes."

THE BANDS - WORDS FIRST, LETTERS RIDE ALONG

BAND	GRADE	RANGE	MEANING
Preferred	A	85–100	Bondable at the best rate.
Standard	B	70–84	Bondable.
Constrained	C	55–69	Bondable with conditions / higher rate.
Decline	D	<55	Not bondable until remediated.

Never render bands as red/green traffic lights - that reads as an eval tool. Tonal copper only.

> **NAMING CONVENTION - RATIFY**
>
> Canonical running-text form: **BackBond** (camel-case), used across every shipped asset. **bond / bonded / bondable** stay lowercase; **BackBond Score** is capitalized. The logo wordmark is set in all-caps as a stylized lockup and does not govern running-text case. **Open item: Eric to ratify camel-case "BackBond" vs. single-capital "Backbond."**

The BackBond Score.

The score is the brand's signature instrument. It is a ratings sheet, not a dashboard: a single 0–100 number on a gauge, four weighted pillars that average to the headline, and a connected A–D band ladder. Reproduce it to spec on every surface.

BACKBOND SCORE
Agent risk rating

SAMPLE

78 /100

RATING BAND
Standard B

0 25 50 75 100

PILLAR · WHAT IT MEASURES	SCORE
Governance Oversight · permissions · audit trail	88
Robustness Prompt injection · jailbreaks · failures	66
Bias & Fairness Disparate treatment · data sensitivity	84
Operations Monitoring · drift · incident response	74

Preferred	Standard	Constrained	Decline
A · 85+	B · 70–84	C · 55–69	D · <55

The band sets whether, and at what rate, the agent is **bondable**. Same agent scores differently at different companies: the score reflects its data access, permissions, and deployment context, not just the model.

How we say it.

TAGLINE OF RECORD

Guardrails lower the odds. We cover the loss.

The line we lead with everywhere. It is the ‘insure’ half of the Focus, and no one who does not pay claims can say it. The logo carries the category; this carries the promise.

THE DIFFERENTIATOR - PROTECT VERBATIM

AI agent insurance: every agent carries its own coverage. Priced per agent, per run, the way AI is billed. That is the ownable mechanic: generic AI-insurance entrants and traditional carriers price entities once a year, while we price the agent by the action. Keep **per agent, per run** intact wherever it appears.

THE LEDE FRAME

Lead where the market actually is - **agents talking to customers** - then escalate to commitments, money, and real actions. Do not open on "moving money" as if every agent wires funds on day one.

THE UNIVERSAL PAIN (CROSS-PERSONA)

"I have agents in production whose decisions can damage my company, my customers, or my career, and I have no defensible measurement of their risk to show my board, my regulator, or my insurer."

THE TAILWIND (FOOTNOTE-LEVEL ONLY)

From January 2026, ISO/Verisk endorsements let carriers exclude generative-AI exposure from standard policies. Real, but a supporting note - never the thesis.

06 · AUDIENCES

Three rooms, one Focus.

PARTNERS

Agent builders & platforms. Win: integrate the rating before their enterprise buyers demand it; coverage built into the sale. Voice: peer-to-peer, technical.

BUYERS

Roles with consequential agent exposure - CFO, COO, Fiduciary. Win: a defensible number and a policy to show the board, regulator, insurer. Voice: capital-fluent, serious.

INVESTORS

AI-infra, fintech, insuretech theses. Win: conviction that this is the category king with a loss-data moat and inevitable demand. Why now: agents are moving from pilots into production, and the rating-and-insurance layer is the gating infrastructure for what comes next. Voice: thesis-grade.

NEAR-TERM PRIORITY

Investors first, then Partners, then Buyers. The immediate sequence starts with capital: Strategic Narrative deck → investor narrative deck. Other tracks follow unless Eric directs otherwise.

LINKAGE - MANDATORY

Every audience hears different proof and different stakes, but a stranger must recover the same Focus from each. Quarterly: read the last ten assets across audiences; if a message has drifted into a smaller or different promise, retire it.

07 · COMPETITIVE

Beat each one on a different axis.

THE ALTERNATIVE	OUR STRUCTURAL AXIS
Vendor self-rating	Third-party + capital-backed. Their score is marketing; ours is bondable.
In-house GRC rubrics	Capital. A rubric produces a number, not a hedge. The board cannot post a spreadsheet against an incident.
AI eval startups	Action-level rating + capital. They evaluate models and ship a PDF; we rate actions and pay the claim.
Hyperscaler bundling	Independence + liability. They cannot credibly rate agents on their own stack, and will not underwrite. Watch for "agent trust" launches; escalate.
Big-four consultancies	Outcomes vs. hours. They sell governance hours; we sell rated outcomes with capital behind them.
Doing nothing	Cost asymmetry. Waiting compounds; a rated, insured deployment is fixed.

THE CARRIER RULE - LOAD-BEARING

Insurance carriers are **partners, not enemies** - future reinsurance and distribution. Never write competitive content against a named carrier. Never claim we "replace" carrier capacity; we provide the methodology and primary layer, they provide the backstop. **Any external asset that names a specific carrier escalates to Eric before shipping.** Violating this is a fireable mistake.

Navy is the institution. Copper is the capital.

Ink Navy
Base, all primary text, dark fields, the score card. Near-black on purpose.

`#0B1A38`

Copper - on dark
The accent on navy: the score, the active band, the payout.

`#D58E5C`

Copper - on white
Same accent, value-tuned for white fields: the wordmark rule, the headline accent word, links.

`#BD6A34`

Slate - secondary
Captions, labels, footnotes. The print-legible floor for small text on white.

`#535E77`

THE COPPER LAW

Copper marks bond, capital, or payout - nothing else. If you cannot name the capital reason a thing is copper, it is navy. Approved moments: the wordmark rule, the headline accent word (bonded / bondable), the score, the active band, a payout figure. **Never** as decoration, never a gradient. Two-tone by background: #D58E5C on navy, #BD6A34 on white. Bands stay tonal, never traffic-light.

Print discipline: light on-screen grays go faint in ink. Secondary text on white never goes lighter than #535E77; the score card's light-on-navy text is print-safe as-is.

09 · TYPE & MARK

The mark is a serif. The collateral is a system.

THE LOGO

BACKBOND
AGENT RISK RATINGS & INSURANCE

Institutional condensed serif wordmark + a 1pt copper underwriting rule + the descriptor **AGENT RISK RATINGS & INSURANCE**. Clear space = the cap-height of "B". Light version on light fields, dark on dark. Do not recolor the rule, add shadow, or stretch. Concept-stage raster - vectorize (.ai/.svg/.eps) before scaled production.

COLLATERAL TYPOGRAPHY

Display	**Archivo ExtraBold**	Headlines, the score
Body	Inter - Regular / Medium / Bold	All running copy
Data	IBM Plex Mono	Figures, scores, ranges

TWO TYPE ROLES, NEVER MIXED UP

The **serif** belongs to the logo only - the institutional signature. **Collateral** (one-pagers, decks, web) is set in Archivo / Inter / IBM Plex Mono. The score card is always a ratings sheet: figures in mono, score on a gauge, bands as one A–D ladder.

10 · VOICE

Operator-grade, or nothing.

CHARACTERISTICS

- **Operator-grade.** Reads like a Chief Risk Officer wrote it, not a content marketer.
- **Specific.** Numbers, names, examples. Vague language is banned.
- **Confident without bravado.** State truth as truth; no "we believe," no bragging.
- **Risk-fluent.** Earn the insurance vocabulary; don't borrow it.
- **Sparse.** No padding, no metaphor that doesn't carry weight.

THE ANTI-TELL CHECKLIST

- No **em-dashes**. Use a regular dash.
- No **rule-of-three** padding where two items do.
- No "in today's [adjective] landscape / era."
- No **empowering, unlocking, seamless, game-changing, revolutionary, paradigm.**
- No "trust" as an undefined noun; no "responsible AI" without a measurable claim.
- No **overclaim**: not "surety," not "guarantee," not "without the downside." We pay the covered loss up to the limit.
- No paragraph a competitor could publish verbatim.

USE LIBERALLY

actuarial · bonded · underwritten · capital-backed · exposure · claims-paid · portfolio · reserves · primary layer · per-agent, per-run · procurement requirement

THE JOURNALIST TEST

Before shipping: would a reporter quote this sentence, or skip it? If every sentence is skippable, the asset is generic. Strengthen the strongest line until it could be the headline.

What's true, and what's locked.

VERIFIED PROOF POINTS

1,180+ probes across 27 risk dimensions	The scan's measured scope. Exact.
Deterministic 0–100 score	The score is deterministic; logs are monitored monthly for drift.
$85B+ annualized quoted premiums, 39 markets	bolttech, public. The founders' prior company.
Ari Katz & Eric Gewirtzman	Founders of Bolt (now bolttech). Insurance operators.

FAL - THE SUPREME LAW, RUN ON EVERY ASSET

Focus: could a stranger recover the Focus from this alone? **Alignment:** does any line imply a different promise? Cut it. **Linkage:** does it share a phrase or proof with other assets? If any fails, it does not ship.

LOCKED

The Focus · the eval-vs-action defense · the naming & scoring system · the visual system · the Carrier Rule.

ESCALATE TO ERIC

Any Focus change · naming a competitor or a carrier · an unverified proof point · a public benchmark we haven't published.

OPEN - AWAITING RATIFICATION

Running-text case: **BackBond** (recommended, matches every shipped asset) vs. "Backbond." One decision, then propagate.

APPENDIX A · COLOR TOKENS

The full palette.

The core three carry the brand; the support tokens keep documents consistent. These are the exact values used in this document. The Copper Law still governs: copper marks bond, capital, or payout only, never decoration; two-tone by background.

TOKEN	HEX	USE
Ink Navy	#0B1A38	Base; all primary text; dark fields; the score card
Copper (on dark)	#D58E5C	Accent on navy: the score, the active band, a payout
Copper Ink (on white)	#BD6A34	Same accent on white: the wordmark rule, the headline accent word, links
Navy-Gray (body)	#36425C	Body copy on white
Slate (secondary)	#535E77	Eyebrows, captions, footnotes. The print-legible floor on white
Card Muted	#8A9BC0	Secondary text on navy
Card Line	#26375C	Dividers on navy
Tint	#F5F7FB	Soft panels on white
Tint Line	#DBE1EE	Panel borders
Rule	#E3E7EE	Hairlines on white
Paper	#FFFFFF	Document field

DO

Spend the accent once: copper = bond / capital / payout. Keep navy near-black and let copper do the work. Sentence-case headlines; uppercase only for tracked eyebrows. Flat color fields; tonal bands within the system.

DON'T

No gradients or blue-purple mesh: the AI-startup tell. No copper as decoration on headers, bullets, or rules. No mid-blue: it reads generic-corporate. No traffic-light bands. No Archivo for body text.

APPENDIX B · TYPE & LOGO

Build specifications.

TYPE RAMP

ROLE	FACE / WEIGHT	SIZE / TRACKING	CASE / LEADING
Display	Archivo 800	33 / -.02em	Sentence · 1.08
Numeral (score)	Archivo 800	46 / -.01em	tabular · 0.9
Title	Archivo 700	18 / -.01em	Sentence · 1.2
Subhead	Inter 700	15	Sentence · 1.35
Body	Inter 400	13	Sentence · 1.55
Emphasis	Inter 700	inline	color Ink Navy
Eyebrow	Inter 700	10 / +.13em	UPPERCASE · 1.0
Caption	Inter 400–500	9–10	Sentence · 1.45
Data	IBM Plex Mono 500	10–12	figures · tables

Sizes are px at one-pager print scale (reference). For slides and web, multiply display sizes ~1.5–2x and set body at 16–18px; keep weights, tracking, and case unchanged. The wordmark is the logo asset (institutional serif), not set in Archivo.

LOGO PRODUCTION SPEC

- **Clear space:** at least the cap-height of the “B” on all sides; keep all other elements out of that zone.
- **Minimum size:** 80px wide digital, 25mm wide print.
- **Background pairing:** light version on light fields, dark version on dark. Never light-on-light or dark-on-dark.
- **Do not** recolor the rule or descriptor, add a shadow, or stretch / condense the wordmark.
- **Descriptor** is set in the wordmark’s serif; match it when vectorizing. Current asset is concept-stage raster.
- **Deliver** a vector master (.ai / .svg / .eps / .pdf) plus transparent PNG at 2x and 3x.
- **Wordmark face:** Frank Ruhl Libre, set large-and-small caps - the serif to match when vectorizing.